

Contents

| | | |
|----------|---|-----------|
| 1 | Introduction | 1 |
| 1.1 | nlmeU R package | 1 |
| 1.2 | ARMD | 1 |
| 2 | Back Ground Materials | 2 |
| 2.1 | Cook's Distance | 2 |
| 3 | Influence Diagnostics for lme objects. | 2 |
| 3.1 | Preparatory Steps | 2 |
| 4 | Influence Diagnostics | 6 |
| 5 | 20.3.2.3. Cook's Distance for the β estimates | 13 |
| 6 | Simulation of the Dependent Variable | 16 |

1 Introduction

In this section we aim to replicate the methods proposed by Galecki et al and transfer to them to the the models fitted by Roy's methodology fitted to the Blood data, as used in Bland and Altman's papers.

1.1 nlmeU R package

This monograph will use functions from the nlmeU package (Galecki and Burzykowski). This package consists of training datasets and utility functions enhancing functionality of nlme package.

logLik1.lme Calculates contribution of one subject to the log-likelihood for lme object

Pwr Calculates power based on a model fit

sigma Extract scale parameter sigma from a model fit

simulateY Simulates values of the dependent variable based on a model fit

1.2 ARMD

Galecki and Burzykowski use in their examples the ARMD data set, sata from Age-Related Macular Degeneration (ARMD) clinical trial. The ARMD data arise from a randomized multi-center clinical trial comparing an experimental treatment (interferon-alpha) versus placebo for patients diagnosed with ARMD.

2 Back Ground Materials

Galecki and Burzykowski revise several core topics before proceeding to introduce their proposed methodology.

2.1 Cook's Distance

3 Influence Diagnostics for lme objects.

- Cook's Distance
- Likelihood Distance

3.1 Preparatory Steps

```
> summary(fm16.5)
Linear mixed-effects model fit by REML
Data: armd
AIC   BIC   logLik
6444.9 6483 -3214.5

Random effects:
Formula: ~time | subject
Structure: Diagonal
(Intercept)      time Residual
StdDev:          7.2357 0.28102   5.0391

Variance function:
Structure: Power of variance covariate
Formula: ~time
Parameter estimates:
power
0.11052
Fixed effects: list(lm3.form)
Value Std.Error  DF  t-value p-value
(Intercept)    5.4416   2.26187 632    2.4058  0.0164
visual0         0.8998   0.03822 231   23.5464  0.0000
time           -0.2416   0.02392 632  -10.0997  0.0000
treat.fActive -2.6553   1.12868 231   -2.3525  0.0195
Correlation:
(Intr) visual0 time
visual0      -0.934
time         -0.071  0.002
treat.fActive -0.270  0.026 -0.002

Standardized Within-Group Residuals:
Min          Q1          Med          Q3          Max
```

-4.148427 -0.329544 0.051823 0.444189 2.930748

Number of Observations: 867

Number of Groups: 234

- The `formula()` can recall the definition of the model defining the mean structure.
- Auxillary Function `logLik1()` which is designed to calculate a contribution of a given subject to the overall likelihood for a given model.
- The number of degrees of freedom reported by `loglik` is equal to 8. This corresponds to the total number of parameters in the model.

Part A

```
<- update(fm16.5)
#####
### code chunk: Chap20.3influence_init
#####

options(width = 65, digits = 5, show.signif.stars = FALSE)
date()
packageVersion("nlmeU")
packageVersion("nlme")
packageVersion("lattice")

sessionInfo()
require(nlme)
require(lattice)

data(armd, package="nlmeU")

## Model M16.5
lm3.form <-                # (12.9, 16.17)
formula(visual ~ visual0 + time + treat.f)
fm16.5 <-                  # R16.13
lme(lm3.form,
random = list(subject = pdDiag(~time)),
weights =varPower(form=~time),
data = armd)

logLik()
```

Part B

```
beta0 <- fixef(fm16.5ml)

beta0

colnames(vcovb) <- names(beta0)
vcovb
```

We extract the β estimates and their estimated variance covariance matrix.

Towards this end, we use the functions `fixef()` and `vcov()` respectively, we can save these estimates and the matrix as objects `beta0` and `vcovb` respectively.

With the help of the `abbreviate` function the names of the beta estimates are shortened to simplify the display of content.

An auxillary function `logLik1()`

we used the `logLik()` function to obtain the value of the log-likelihood for the fitted model. It should be noted that the function returns the log-likelihood evaluated at the set of the

The auxillary functon `logLik1()` has been included in the package `nlmeU`.

estimated fixed effects and variance-covariance parameters and for the data set, to which the model is fitted. In the context of influence diagnostics, we need a more general function that allows

modfit an object of class *lme* representing the lme model fitted to a given dataset using ML estimation.

dt1 a dataframe with data for one subject, for whom the likelihood function to be evaluated

dtInit an optional auxillary data frame

The dataframe provided in the argument `dt1` is typically created by choosing a subset with one subject from the data used to obtain the model fit object specified in the `modfit` argument. However, in general, any plausible data for one subject, not necessarily from the dataset used to fit the model, can be used.

The auxillary data provided in the argument `dtInit` is temporarily appended to the `dt1` data. during the `logLik()` function execution.

The `logLik1()` function returns the numeric contributions of the single subject, with the data specified in the `dt1` argument, to the log likelihood for the model specified in the `modfit` argument.

Contributions of Individual Subjects to the log-likelihood for fitted model

```
lLik.i <- as.vector(lLik.i)

lLik.i[1:5]
sum(lLik.i)
```

Plot of individual contributions to the log-likelihood (traditional graphics)

```
subject.c <-
subject.x <-
plot()
points()
text()
```

next we use the function `logLik1()` to compute the loglikelihood contributions for all subjects.

```
#####
### code chunk: R20.7a
#####
require(nlmeU)
df1 <- subset(armd, subject %in% "1")          # Data for subject "1"
logLik1(fm16.5ml, df1)                        # logLik_i for subject "1"

lLik.i <- by(armd, armd$subject,
FUN = function(dfi) logLik1(fm16.5ml, dfi))
lLik.i <- as.vector(lLik.i)    # Coerse array to vector
lLik.i[1:5]                    # logLik_i for the first five subjects
sum(lLik.i)                    # Sum logLik_i; compare to Panel 20.6a

#####
### code chunk: R20.7b
#####
nx <- by(armd, armd$subject, nrow)            # ni
lLik.n <- lLik.i/as.vector(nx)                # logLik_i
outL <- lLik.n < -6                           # TRUE for values < -6
lLik.n[outL]                                  # logLik_i/ni < -6
subject.c <- levels(armd$subject)
subject.x <- as.numeric(subject.c)

plot(lLik.n ~ subject.x, type = "h")           # Fig. 20.1
points(subject.x[outL], lLik.n[outL], type = "p", pch = 16)
text(subject.x[outL], lLik.n[outL], subject.c[outL])

#####
### code chunk: R20.8a
#####
lmeU <- function(cx) {
dfU <- subset(armd, subject != cx)            # L00 data
update(fm16.5ml, data = dfU)                  # L00 fit
}

lmeUall <- lapply(subject.c, lmeU)             # List with L00 fits
names(lmeUall) <- subject.c                   # Names assigned
```

We present the syntax to plot the per-observation individual log-likelihood contributions. First, with the help of the `by()` function, we create the array `nx`, which contains the number of observations

4 Influence Diagnostics

We use the results of the preparatory steps to perform influence-diagnostic calculations for the model. More specifically we evaluate the influence of every subject included in the data set.

We create a list containing the results of fitted model the “leave-one-subject-out” (LOO) datasets and explore its contents.

We define the function `lmeU()`, which fits the model to the data from the `armd` when the function `lmeU()` is executed, and LOO data frame, named `dfU`, is created with the subject, indicated by the `cx` argument.

Subsequently model is fitted by `dfI` by applying the next, with the help of the function `lapply()`, we apply the `lmeU()` to the consecutive elements of the character vector `subject.c`.

As a result, we obtain the list `lmeUall`, with `lme`-class model fit objects as elements. The model-fit objects contains the result of fitting model M16.5.

Finally, we name the components of the `lmeUall()` list using the subjects identifier stored in the vector `subject.c`.

This technique is computationally expensive, as it required the model to be fitted m number of times, omitting one of the m cases each time.

Execution time can be improved if we decided to perform a reduced number of likelihood iterations, instead of performing iterations until there is convergence.

The values, based on the first few iterations, are expected to give a fairly good approximation of the LOO estimates.

The names of the first six components are printed out using the function `names()`.

To extract the LOO data frame for, e.g., the subject “6”, we refer to the “6” component of the `lmeUall` list.

The extracted data frame is stored in the object `dataU6`. By using the function `dim()` we can check the dimensions of the data frame.

Model is fitted to a sequence of “leave-one-subject-out” out data sets.

```
# back - what is cx?
# cx is Case identity
lmeU <- function(cx){
dfU <- subset(myData, subject !=cx)
update(mymodel,data=dfU)
}

# what is lmeU?

lmeUall <- lapply(subject.c, lmeU)
```


Exploring the contents of the lmeUall object.

```
names(lmeUall)
dataU6 <- lmeUall[["6"]]$data
dim(dataU6)
unique(dataU6$subject)[1:6]
```

```
#####
### code chunk: R20.6a
#####
fm16.5ml <- update(fm16.5, method = "ML") # ML estimation
formula(fm16.5ml)                        # Recall model formula.
fm16.5ml$call$data                       # Recall data name.
logLik(fm16.5ml)                         # Log-likelihood value

#####
### code chunk: R20.6b
#####
beta0 <- fixef(fm16.5ml)                  # beta
names(beta0)                             # Long names
names(beta0) <- abbreviate(names(beta0), minlength = 7) # Short names
beta0                                     # beta printed.
vcovb <- vcov(fm16.5ml)                   # vcovb
colnames(vcovb) <- names(beta0)           # Short names
vcovb                                     # vcovb printed.
```

Galecki presents the code used to calculate and plot individual likelihood displacements.

For an LMM, it is required the computation of the full log-likelihood for $\hat{\theta}$, the ML estimate for θ obtained by fitting the model to all data, and for $\hat{\theta}_{(-i)}$, the ML estimate obtained by fitting the model to the data with the i -th subject excluded.

Note that both values of the log-likelihood, used in the definition of the likelihood displacement, should be calculated taking into account all observations.

Galecki creates an auxiliary function `lLik()` which, for a given subject indicated by the main argument, extracts the lme model fit object for the corresponding LOO data.

The corresponding log-likelihood function is extracted from the lmeU with the help of the `logLik()`

The returned value `lLikU + lLik.s` is the log-likelihood evaluated for all observations, using the displaced estimates of the model parameters.

Calculation of the likelihood displacement

```
lLik <- function(cx)
{
  lmeU <- lmeUall[[cx]]
  lLikU <- logLik(lmeU, REML = FALSE)
  df.s <- subset(armd, subject == cx)

  lLik.s <- logLik1(lmeU, df.s)

  return(lLikU + lLik.s)
}

lLikUall <- apply(subject.c, lLik)
```

```
#####
### code chunk: R20.8b
#####
names(lmeUall)[1:6]
dataU6 <- lmeUall[["6"]]$data      # LOO data for subject 6
dim(dataU6)                       # Number of rows is 863
unique(dataU6$subject)[1:6]        # Subject no. 6 omitted

#####
### code chunk: R20.9a
#####
```

```

lLik <- function(cx){
lmeU   <- lmeUall[[cx]]           # LOO fit extracted
lLikU  <- logLik(lmeU, REML = FALSE) # LOO log-likelihood
df.s   <-                        # Data for subject cx...
subset(armd, subject == cx)
lLik.s <- logLik1(lmeU, df.s)      # ...and log-likelihood.
return(lLikU + lLik.s)           # "Displaced" log-likelihood...
}
lLikUall <- sapply(subject.c, lLik) # ...for all subjects.

dif.2Lik <- 2*(logLik(fm16.5ml) - lLikUall) # Vector of LDi
summary(dif.2Lik)

```

```

names(dif.2Lik) <- subject.c #subjects ids assigned
outL <- dif.2Lik > 0.5

dif.2Lik[outL]

## Plot Component

library(lattice)

subject.f <- factor(subject.c, levels = subject.c)

myPanel <- function(x,y, ...){
  x1 <- as.numeric(x)
  panel.xyplot(x1,y, ... )
  ltext(x1[outL],y[outL], subject.c[outL]) # outlying LDis
}

dtp <- dotplot(dif.2Lik ~subject.f, panel = myPanel, type= "h")

# ggplot?
lxlims <- length(dtp$x.limits)

update(dtp,xlim=rep(" ", lxlims),grid= "h")

```

By applying the `summary()` function to the vector, we obtain summary statistics of the computed likelihood-displacement values.

we create the logical output vector `outL` which indicates the subjects with the values of the likelihood displacement exceeding say 0.5.

From the printout of the selected elements of the vector `dif.2Lik` it follows that there are seven such subjects.

We then use the function `dotplot()` from the package `lattice` to plot the likelihood distance value for all subjects.

The x-axis of the plot is constructed using numeric representation of the `numeric.f` factor, containing values ranging from 1 to 234.

```
#####
### code chunk: R20.9b
#####
names(dif.2Lik) <- subject.c          # Subjects' ids assigned
outL <- dif.2Lik > 0.5                # Outlying LDi's
dif.2Lik[outL]
library(lattice)
subject.f <- factor(subject.c, levels = subject.c)
myPanel <- function(x, y, ...){
  x1 <- as.numeric(x)
  panel.xyplot(x1, y, ...)
  ltext(x1[outL], y[outL], subject.c[outL]) # Label outlying LDi
}

dtp <-                                # Fig. 20.2
dotplot(dif.2Lik ~ subject.f, panel = myPanel, type = "h")
lxlims <- length(dtp$x.limits)
update(dtp, xlim = rep("", lxlims), grid = "h")
```

5 20.3.2.3. Cook's Distance for the β estimates

Cook's distance for the β estimates was defined in (4.26) for the classical LM. The definition can be extended to the LMMs in a straightforward manner.

Part 2 plot of cook's distance using traditional graphics

```
#####
### code chunk: R20.10a
#####
betaUall <- sapply(lmeUall, fixef)      # Matrix with beta(-i)
vb.inv <- solve(vcovb)
CookDfun <- function(betaU){
  dbetaU <- betaU - beta0                # beta(-i) - beta
  CookD.value <- t(dbetaU) %*% vb.inv %*% dbetaU
}
```

```

CookD.num <- apply(betaUall, 2, CookDfun)
(n.fixeff <- length(beta0))           # Number of fixed effects
rankX <- n.fixeff                     # Rank of matrix X
CookD <- CookD.num/rankX              # Cook's distance Di

```

```

### Blood Data Equivalent
betaUall <- sapply(lmeUall,fixef)

vb.inv <- solve(vcovb)

cookDfun <- function(betaU){
  dbetaU <- betaU - beta0
  cookD.value <- t(dbetaU) %*% vb.inv %*% dbetaU
}

```

Plot of Cook's Distance using traditional graphics. Outlying values annotated.

```

#####
### code chunk: R20.10b
#####
outD <- CookD > 0.03                # Outlying Di's
subject.c[outD]                     # Subjects' ids

plot(CookD ~ subject.c,
     ylab = "Cook's D", type = "h") # Fig. 20.3
text(as.numeric(subject.c[outD]),
     CookD[outD], subject.c[outD])  # Annotation
points(subject.c[outD], CookD[outD])

```

```

### Blood Data Equivalent
outD <- cookD > 0.03

#Create the Stick plot
plot(CookD ~ subject.c,
     ylab="Cook's Distance", type="h")

```

```
text()
points()
```

Cook's distance for the *beta* estimates were defined by 4.26 for the classical LM model.

Next we compute the inverse of the variance covariance matrix $\hat{\beta}$ using the `solve()` function. We store the resulting matrix in the object `vb.inv`.

Subsequently we define the function `cookDfun()` which, for a vector given in the `betaU` argument, computes the value of the numerator of Cook's Distance, as in (4.26).

The function is then applied sequentially to all columns of the matrix `betaUall`.

The resultant vector is divided by the number of the fixed effects coefficients, which, under the assumption that the design matrix is of full rank, is equivalent to the rank of the design matrix.

The outcome is stored in the vector `cookD` and contains the values of cooks distanced for all subjects.

We create the logical vector `outD` which indicates the subjects with cook's distance values that exceed 0.03.

Present the scatterplot matrix of the two-dimensional projections of the differences. for all pairs of the fixed-effects coefficients.

The plot was generated using the `spiom` function. The main argument of the function was obtained by subtracting the `beta0` vector from the rows of the transposed `betaUall` matrix.

The labels used in the panels located on the diagonal of the figure provide the fixed estimates of fixed effects coefficients of the model and their estimates SEs. The panels above the diagonal include points for all subjects. The points for non-outlying values are plotted using small size open circles. The five outlying values are plotted using different plotting symbols defined in the legend of the figure at the top of the graph.

The effect of removal of this subject on the estimates of the remaining fixed-effects coefficients is relatively small. In contrast, removing the subject "227" affects the estimate of all fixed-effects coefficients to a different degree and in different directions.

More specifically, the intercept is driven towards lower values : the positive slope associated with visual acuity at baseline `visual0`, is further increased, the negative slope associated with `time` is brought closer to zero, and the treatment effect is attenuated.

Overall we note that the effect of removing any of the subjects on the fixed effects estimates is very small, as it does not exceed 0.05 of the SE of any of the estimates.

6 Simulation of the Dependent Variable

```
### Note: Simulations in Panel R20.11 take a long time
```

```
#####
```

```
### code chunk: Chap20.4init
```

```
#####
```

```
options(width = 65, digits = 5, show.signif.stars = FALSE)
date()
```

```
sessionInfo()
```

```
require(nlme)      #
```

```
data(armd, package="nlmeU")
```

```
lm3.form    <-                                # (12.9)
formula(visual ~ visual0 + time + treat.f)
fm16.5    <-      #update(fm16.4,              # M16.5 <- M16.4
lme (lm3.form,
random = list(subject = pdDiag(~time)),
weights = varPower(form = ~time), #
data = armd)
fm16.5ml <- update(fm16.5, method = "ML")
```

```
#####
```

```
### code chunk: R20.11
```

```
#####
```

```
library(nlmeU)
simY <- simulateY(fm16.5ml, nsim = 1000, seed = 1238917) # Simulated y from M16.1
auxDt <- subset(armd,                                     # Auxiliary data
select = c(subject, visual, visual0, time, treat.f))
simYsumm <-
apply(simY,
MARGIN = 2,                                             # Over columns
FUN = function(y){
auxDt$visual <- y                                     # Dependent variable updated
auxFit <-                                             # Update M16.1 with new y
update(fm16.5ml, data = auxDt)
summ <- summary(auxFit)                               # Summary
beta <- fixef(summ)
list(beta = beta)
})
simYsumm[[1]]                                         # beta for the 1st simulation
```



```
#####
### code chunk: R20.12
#####
betaE <- sapply(simYsumm,          # Matrix with betas
FUN = function(x) x$beta)
rowMeans(betaE)                   # Empirical beta (see Panel R20.6b)
cov(t(betaE))                     # Empirical Var(beta)

## sessionInfo
sessionInfo()
```

We consider simulations of the dependent variable, based on the marginal distribution implied by the fitted model. Towards this end, we have developed `simulateY()`, which can be used for objects of class `lme`.

We note that the function is different from `simulate.lme()`, available in the `nlme` in that the latter returns simulation-based REML and/or ML values and not the value of the dependent variable.

We demonstrate the use of the `simulateY()` function to create the empirical distribution of the β estimates. As an example, we consider the model which was fitted to the `armd` data.

Note however that the presented syntax is fairly general and can be used for other LMEs as well.

We apply the function `simulateY()` to the object `fm16.5ml`.

```
library(nlmeU)
simY <- simulateY(fm16.5ml, nsim=1000)

simYsumm <- apply( simY, MARGIN = 2,
FUN = function(y){

  }

)

simYsumm[[1]]

# \hat{\beta} for the 1st simulation
```

There is a second argument `nsim` that indicated how many simulations

The auxiliary function performs the following steps

- the dependent variable `visual` in the `auxDt` data frame is replaced with a new set of simulated values contained in the vector `y`.
- Model `M16.5` is fitted to the modified data frame

- the vector `beta` with the estimates of the fixed effects coefficients is extracted from the summary of the model-fit object with the help of the `fixef()` function.
- The vector of estimates is returned as a list with one component named `beta`.

They are drawn from the marginal distribution of the dependent variable implied by the fitted model.

It should be mentioned that the creation of the `simYsumm` list involves refitting the model many times, and it therefore takes a long time.

Toward this end, with the help of the `sapply()` function, we extract the vectors with the values of $\hat{\beta}$ for each simulation for the list-object `simYsumm` and bind them column-wise into the `betaE`.

Then we use the function `rowMeans()` to compute the mean values of the columns (i.e. across the rows) of `betaE`.